

Parallelized Stochastic Gradient Descent

第八組 張家銘 繆穩慶 麥佩琦



Outline

- Introduction
- Problem statement
- Proposed approach
- Experiment and evaluation
- Conclusion

Introduction


| Gradient Descent

- Iteratively moves toward parameter to find one that **minimize the loss function**

$$x_{t+1} = x_t - \eta g_{GD}$$

Introduction

| Gradient Descent



Introduction

| Gradient Descent

- For each iteration, compute **all pair of examples** in the training set to get gradient

→ **Long computation time**

Introduction


| Stochastic Gradient Descent

- At each iteration, randomly choose **one sample** to calculate
 - Mini-batch SGD: choose a **subset** of training set
- **Much faster**

Problem statement

However, GD/SGD are inherently serial due to their iterative nature

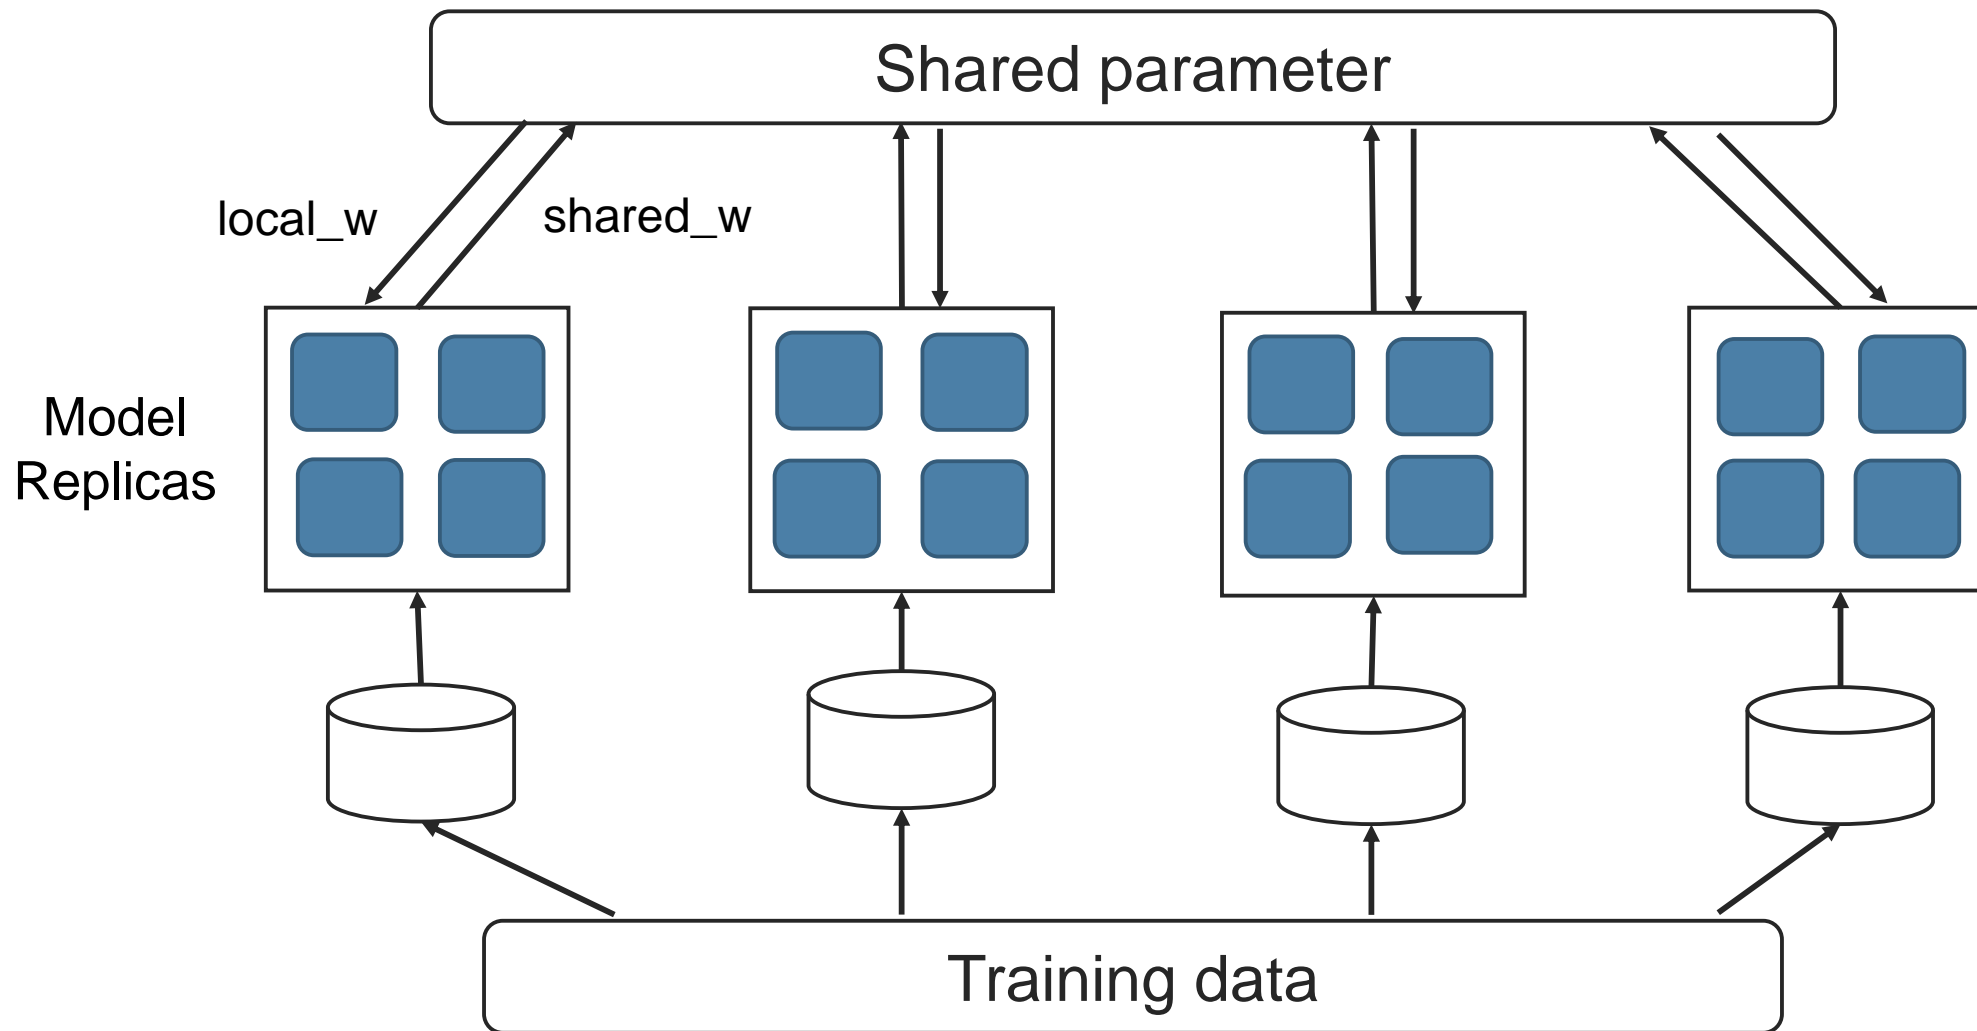
- Do parallelism when mini-batches calculate gradient
- Update parameters with **reduce** technique

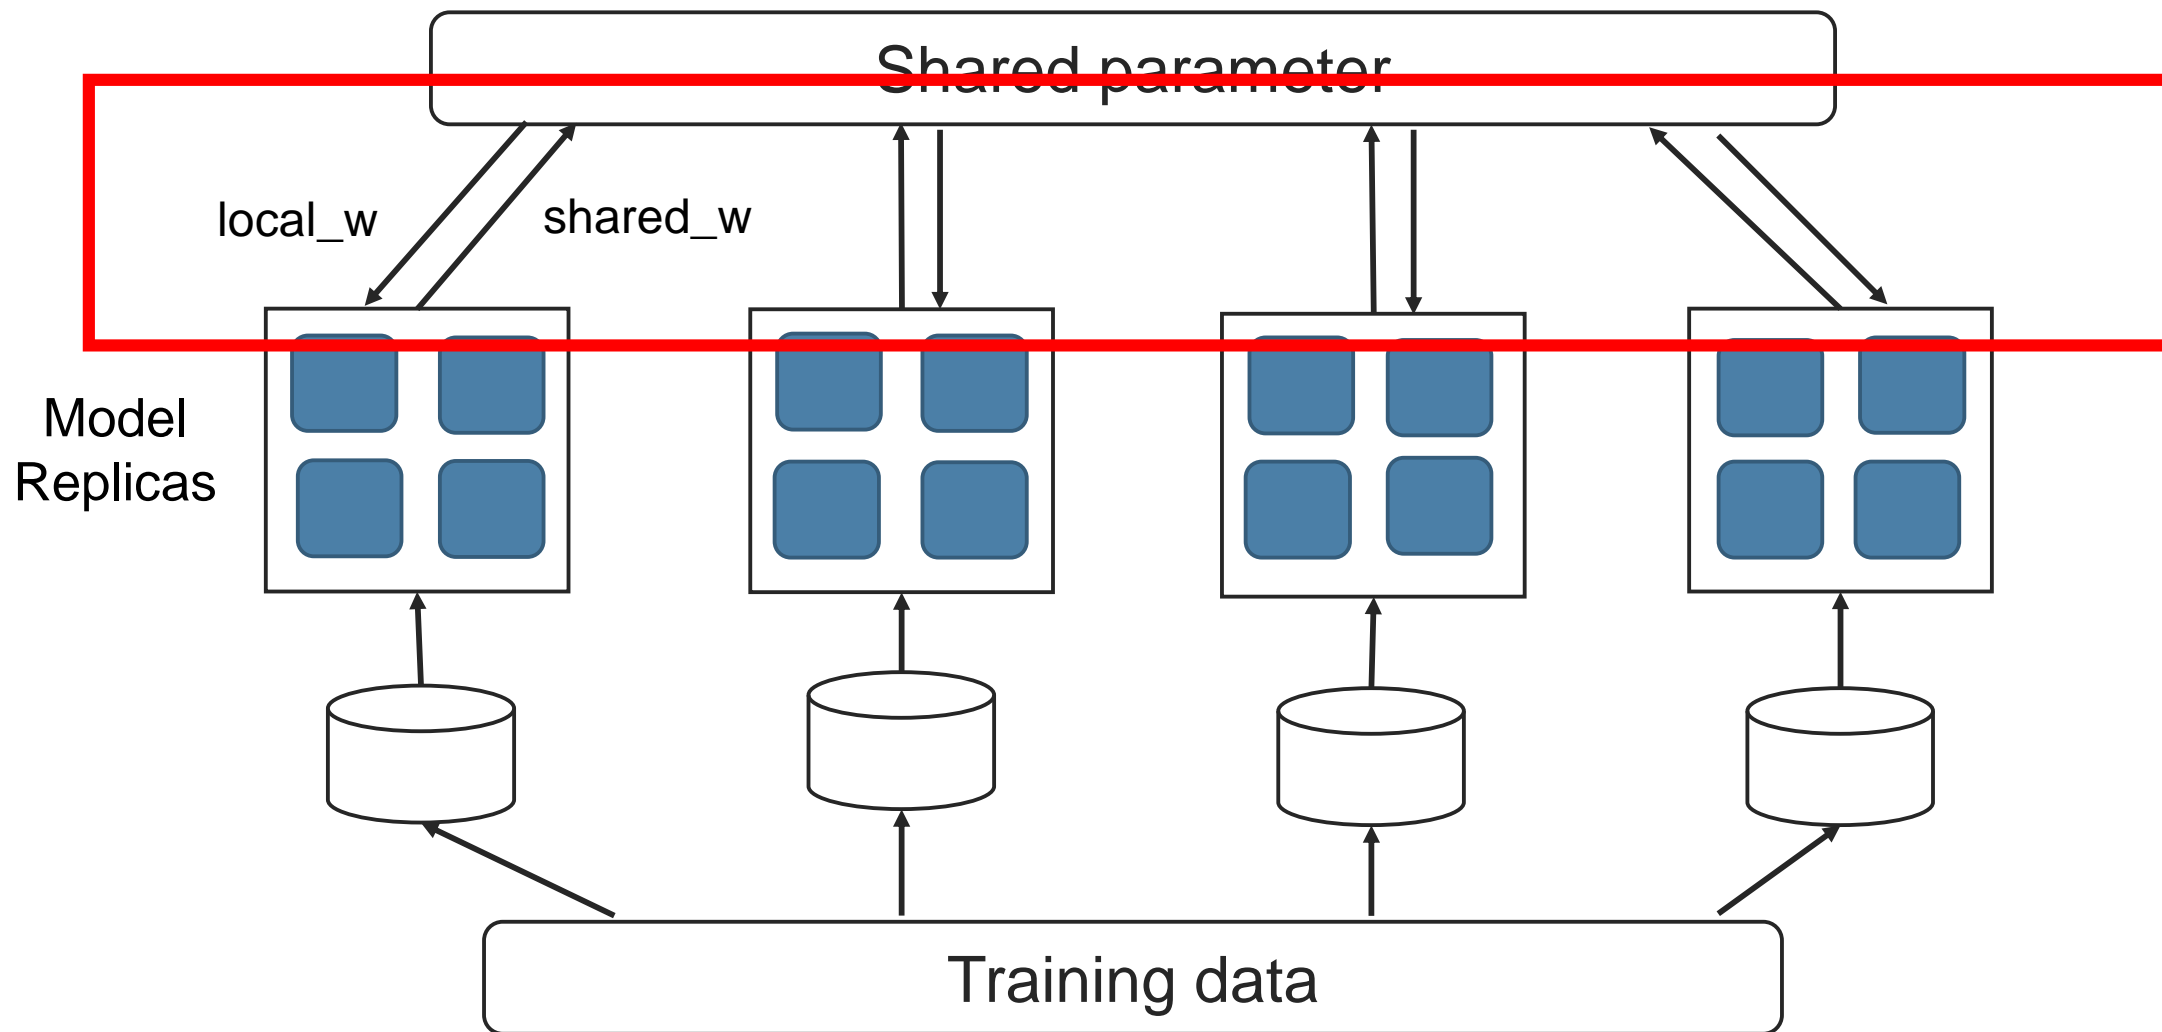


Proposed approach

- Divide training set into several mini-batches. Each mini-batch is assigned to one processor.
- We apply two methods to update the shared parameters.





—





Method 1: Synchronous

- For each processor i
 - Initialize local parameter $w_i = \text{shared parameter } v$
 - Calculate SGD
- With **reduce** technique, update share parameter $v = \frac{1}{c} \sum_{i=1}^c w_i$



Gather

Every Batch

Delta W0

Delta W1

Delta W2

Delta W3


W1 = W

W2 = W

W3 = W

$$W = W - \frac{1}{4} \sum_{i=0}^3 \text{Delta } W_i$$


Broadcast




Method 2: Asynchronous

Prerequisite: the speed of each processor is similar

- For each processor i
 - Calculate SGD with local parameter w_i
- Send w_i to update share parameter v **in turn** and receive updated v as w_i




CORE0



CORE1



CORE2



CORE3

Batch = 0 $W \mathrel{:=} \Delta W$

$W1 \mathrel{:=} \Delta W1$

$W2 \mathrel{:=} \Delta W2$

$W3 \mathrel{:=} \Delta W3$

Batch = 1 $W \mathrel{:=} (\Delta W + \Delta W1)/2$ $W1 = W$

$W2 \mathrel{+}= \Delta W2$

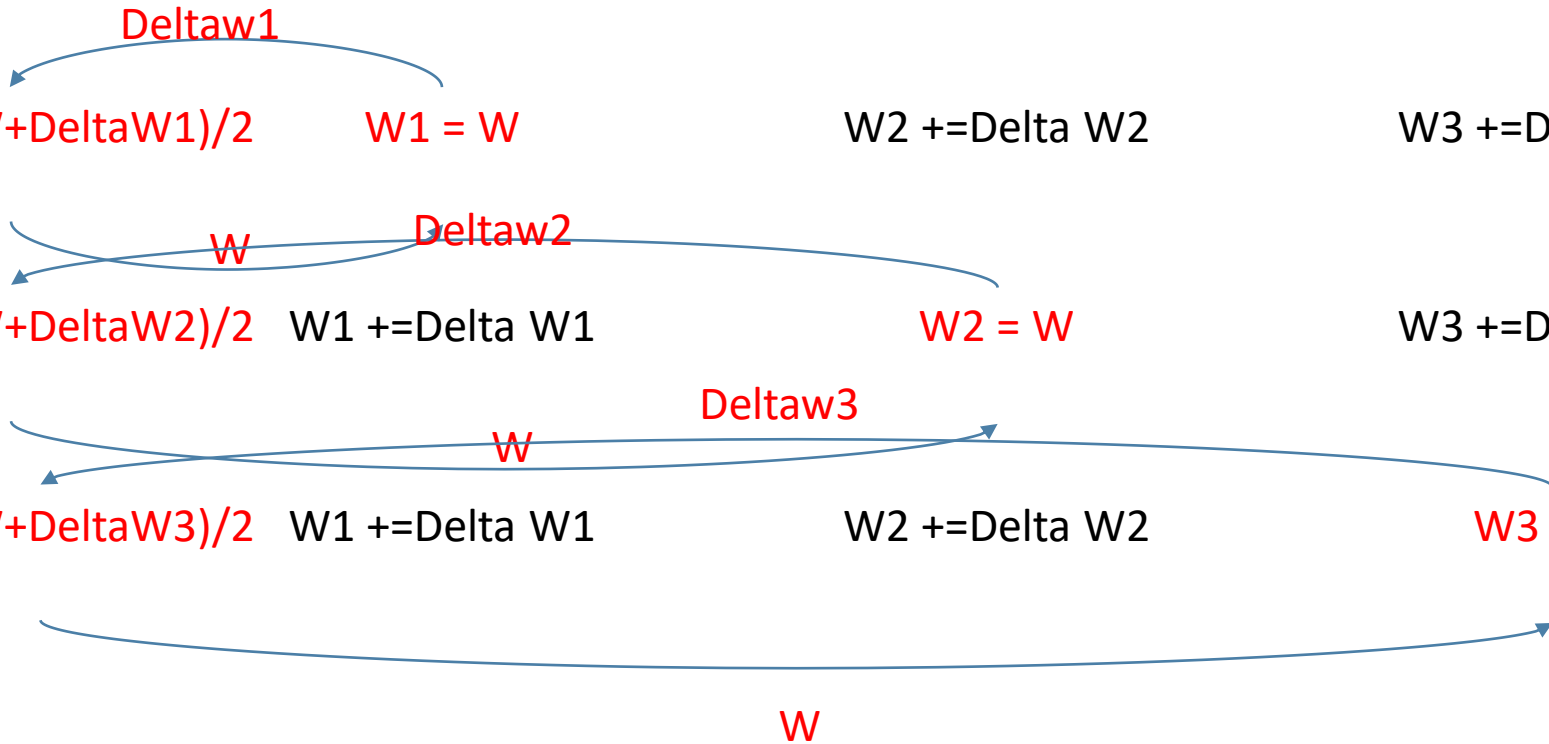
$W3 \mathrel{+}= \Delta W3$


Batch = 2 $W \mathrel{:=} (\Delta W + \Delta W2)/2$ $W1 \mathrel{+}= \Delta W1$

$W2 = W$

$W3 \mathrel{+}= \Delta W3$


Batch = 3 $W \mathrel{:=} (\Delta W + \Delta W3)/2$ $W1 \mathrel{+}= \Delta W1$ $W2 \mathrel{+}= \Delta W2$ $W3 = W$





Experiment & evaluation

Experiment



Hyperparameters :

- Activation function : *sigmoid*
- Learning Rate : 1 (Reduce *lr* by 1% after each epoch)
- Epochs : 5
- Batch size : 100

Optimizer : **mini-batch SGD**

Experiment

Environment :

CPU: Intel i7-8700 3.20 GHz
system memory: 16 GB

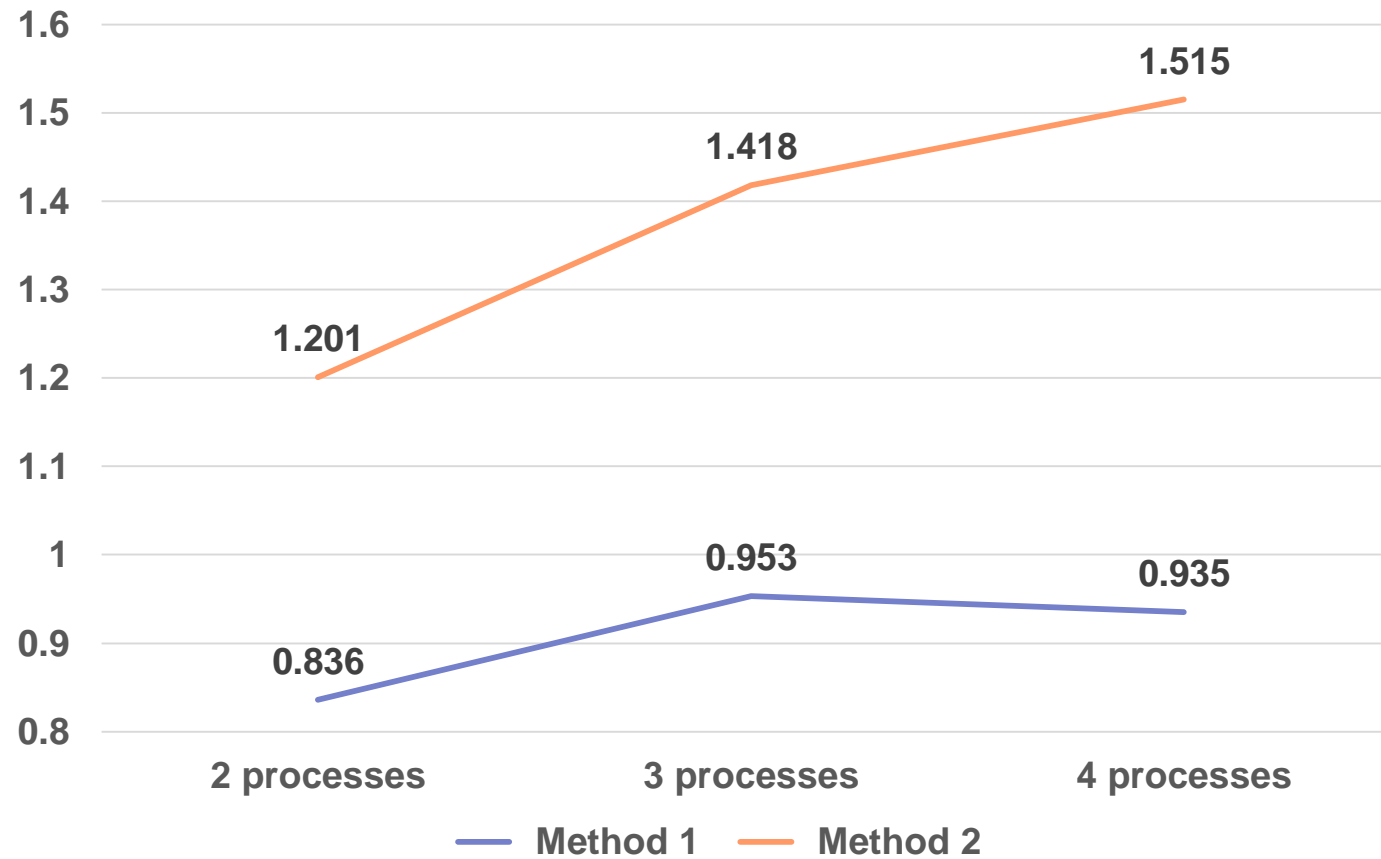
Language : python with mpi4py



Evaluation


Method	# Process	Loss	Accuracy	Time (s)
Serial SGD	1	0.3726	0.8880	9.0821
Sync	2	0.2709	0.9210	10.8651
	3	0.2744	0.9213	9.5284
	4	0.2977	0.9185	9.7092
Async	2	0.3204	0.9111	7.5638
	3	0.6277	0.8426	6.4041
	4	0.4981	0.8596	5.9964

Evaluation




Evaluation

Serial SGD




Synchronous method




Evaluation

Serial SGD



Asynchronous method



Evaluation

Method	# Process	CPU usage	Memory usage
Serial SGD	1	35~40%	12%
Sync	2	40~45%	28.7%
	3	55~57%	36.5%
	4	65~70%	43.3%
Async	2	52~55%	28%
	3	70~75%	35%
	4	80~85%	42.9%

Conclusion

- Provide two methods to parallelize SGD
- The performance of first method is worse than serial program, we think this is because of that communication overhead is too high.
- The second method may cause the decrease of accuracy, this is because of the share weights are not update in some process in each iteration.

Q & A



Team 8